Develop mapreduce progrum to calculate the frequency of a given file.

mapper.py

```
#!/usr/bin/env python3
import sys

# Read input line by line
for line in sys.stdin:
    line = line.strip()  # Remove whitespace
    words = line.split()  # Split into words
    for word in words:
        print(f"{word}\t1")  # Emit (word, 1)
```

reducer.py

```
#!/usr/bin/env python3
import sys
from collections import defaultdict

word_count = defaultdict(int)

# Read input from standard input
for line in sys.stdin:
    word, count = line.strip().split("\t")
    word_count[word] += int(count)

# Print the final counts
for word, count in word_count.items():
    print(f"{word}\t{count}")
```

Output :-

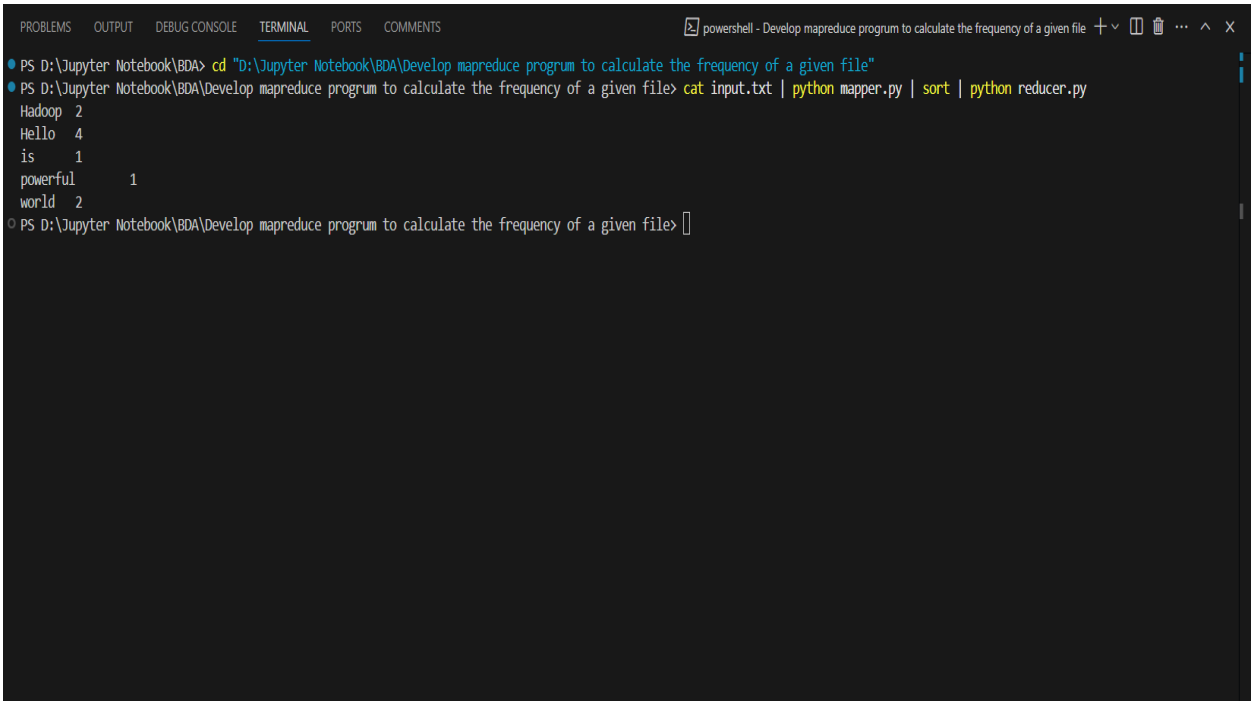
PROBLEMS
OUTPUT
DEBUG CONSOLE
TERMINAL
PORTS
COMMENTS
powershell - Develop mapreduce program to calculate the frequency of a given file
PS D:\Jupyter Notebook\BDA> cd "D:\Jupyter Notebook\BDA\Develop mapreduce program to calculate the frequency of a given file"
PS D:\Jupyter Notebook\BDA\Develop mapreduce program to calculate the frequency of a given file> cat input.txt | python mapper.py | sort | python reducer.py
Hadoop 2
Hello 4
is 1
powerful 1
world 2
PS D:\Jupyter Notebook\BDA\Develop mapreduce program to calculate the frequency of a given file>

Implement matrix multiplication using map-reduce.

matrix.txt

```
A 0 0 1
A 0 1 2
A 1 0 3
A 1 1 4
B 0 0 5
B 0 1 6
B 1 0 7
B 1 1 8
```

mapper.py

```
#!/usr/bin/env python3
import sys

# Read input line by line
for line in sys.stdin:
    line = line.strip()
    matrix, row, col, value = line.split()
    row, col, value = int(row), int(col), int(value)

    if matrix == 'A':  # Emit values to be multiplied with B
        for k in range(2):  # Assuming B has 2 columns
            print(f"{row},{k}\tA,{col},{value}")

    elif matrix == 'B':  # Emit values to be multiplied with A
        for i in range(2):  # Assuming A has 2 rows
            print(f"{i},{col}\tB,{row},{value}")
```

reducer.py

```
#!/usr/bin/env python3
import sys
from collections import defaultdict

# Store values by keys
product_terms = defaultdict(list)

# Read input from standard input
for line in sys.stdin:
    key, value = line.strip().split("\t")
    product_terms[key].append(value)

# Compute matrix multiplication result
result = defaultdict(int)
for key, values in product_terms.items():
    a_values = {int(v.split(",")[1]): int(v.split(",")[2]) for v in values if v.startswith("A")}
    b_values = {int(v.split(",")[1]): int(v.split(",")[2]) for v in values if v.startswith("B")}

    for k in a_values:
        if k in b_values:
            result[key] += a_values[k] * b_values[k]

# Print the final matrix product
for key, value in sorted(result.items()):
    print(f"{key}\t{value}")
```

Output:-

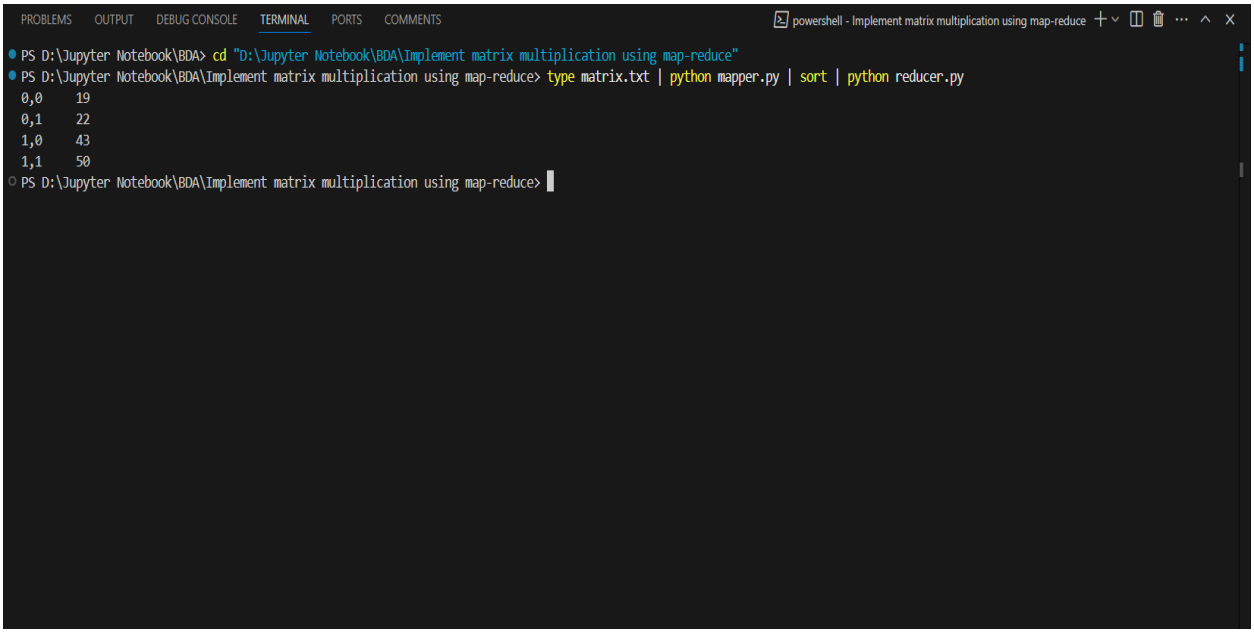
PROBLEMS
OUTPUT
DEBUG CONSOLE
TERMINAL
PORTS
COMMENTS
powershell - Implement matrix multiplication using map-reduce
PS D:\Jupyter Notebook\BDA> cd "D:\Jupyter Notebook\BDA\Implement matrix multiplication using map-reduce"
PS D:\Jupyter Notebook\BDA\Implement matrix multiplication using map-reduce> type matrix.txt | python mapper.py | sort | python reducer.py
0,0 19
0,1 22
1,0 43
1,1 50
PS D:\Jupyter Notebook\BDA\Implement matrix multiplication using map-reduce>

Mongodb Installation 4 creation of database Collection Insert Query, update query & delete Documents.

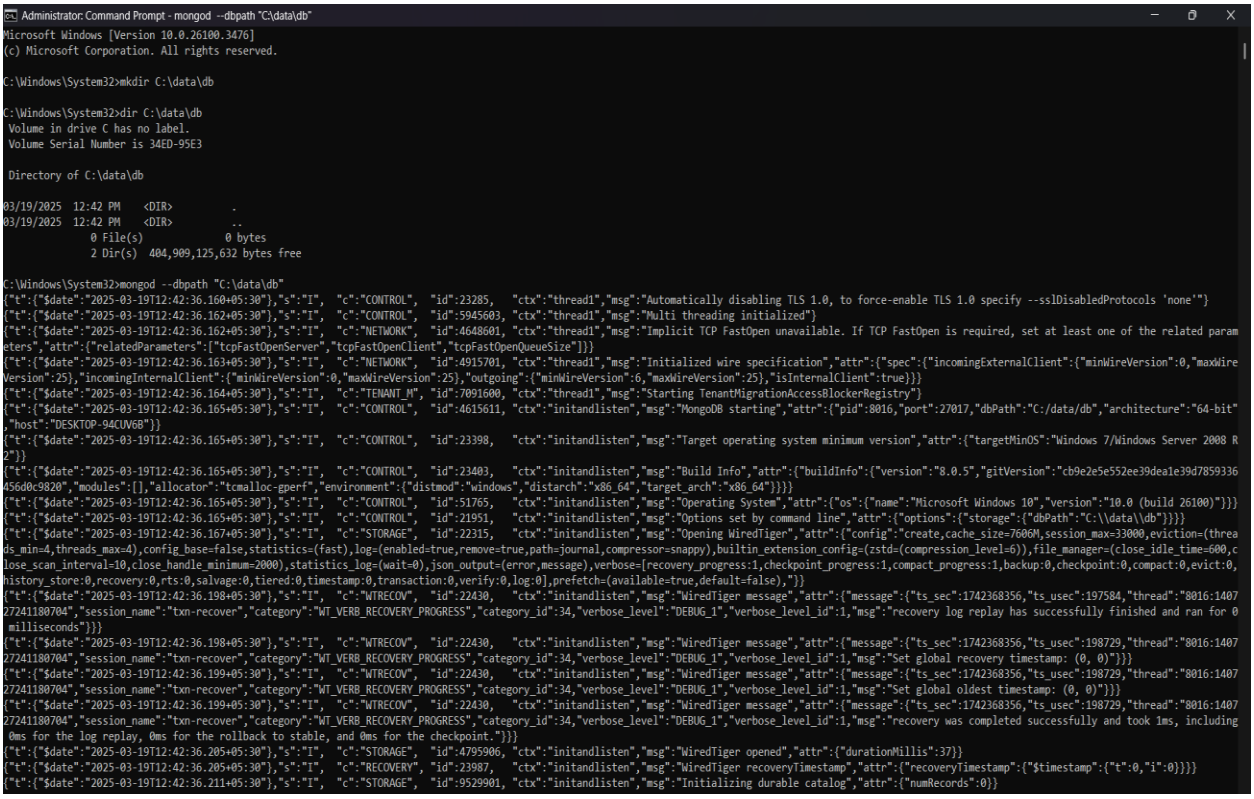

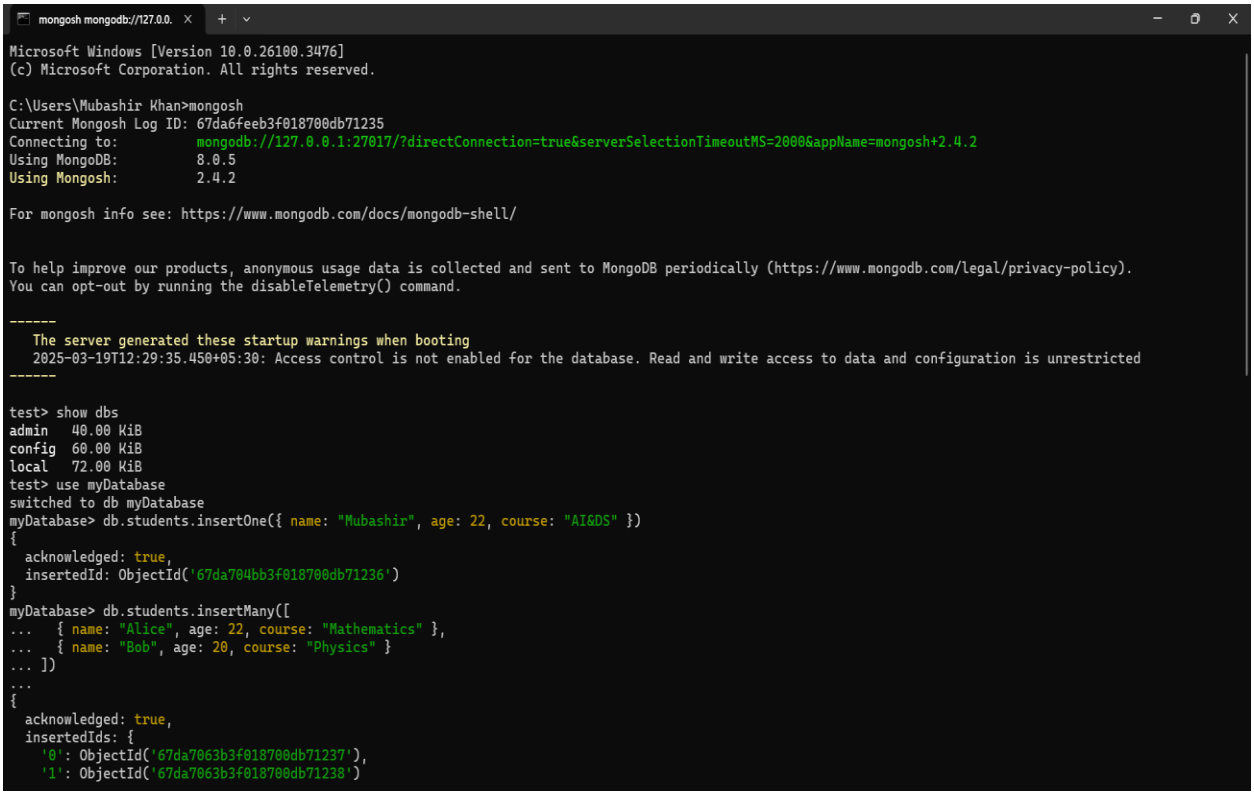

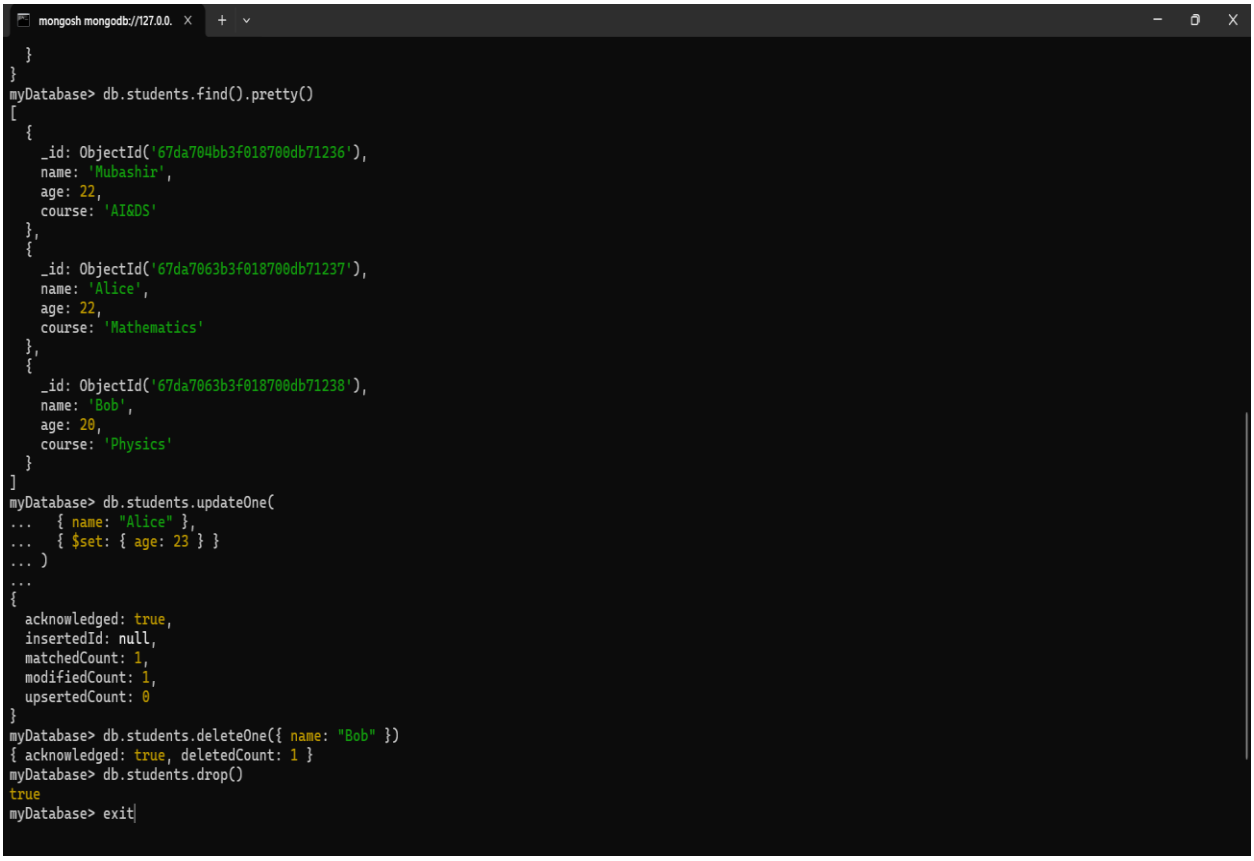
mongosh mongodb://127.0.0.
}
}
myDatabase> db.students.find().pretty()
[
{
_id: ObjectId('67da704bb3f018700db71236'),
name: 'Mubashir',
age: 22,
course: 'AI&DS'
},
{
_id: ObjectId('67da7063b3f018700db71237'),
name: 'Alice',
age: 22,
course: 'Mathematics'
},
{
_id: ObjectId('67da7063b3f018700db71238'),
name: 'Bob',
age: 20,
course: 'Physics'
}
]
myDatabase> db.students.updateOne(
... { name: "Alice" },
... { $set: { age: 23 } }
...)
...
{
acknowledged: true,
insertedId: null,
matchedCount: 1,
modifiedCount: 1,
upsertedCount: 0
}
myDatabase> db.students.deleteOne({ name: "Bob" })
{ acknowledged: true, deletedCount: 1 }
myDatabase> db.students.drop()
true
myDatabase> exit

Visualization connect to data build charts & analyze data create dashboard

dashboard.py

```
import streamlit as st

import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

import plotly.express as px

# 📌 Load Titanic dataset from a public URL

DATA_URL = "https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv"

data = pd.read_csv(DATA_URL)

# ------------------ □ STREAMLIT DASHBOARD START ------------------

# 📌 Set Power BI-like page layout

st.set_page_config(page_title="🚢 Titanic Analysis", layout="wide")

# 📌 Dashboard Header

st.markdown("<h1 style='text-align: center;'>🚢 Titanic Passenger Dashboard</h1>", unsafe_allow_html=True)

st.markdown("<h3 style='text-align: center;'>Analyze survival rates and passenger demographics</h3>", unsafe_allow_html=True)

# ---------------- □ KPI METRICS (TOP ROW) -----------------

st.markdown("### 🚀 Key Statistics")

col1, col2, col3, col4 = st.columns(4)

# 📌 Display KPIs

with col1:

   st.metric(label="👥 Total Passengers", value=f"{len(data)}")

with col2:

   survived_count = data["Survived"].sum()

   st.metric(label="❤️ Survived", value=f"{survived_count}")

with col3:

   avg_fare = round(data["Fare"].mean(), 2)

   st.metric(label="💰 Avg Fare", value=f"${avg_fare}")

with col4:

   avg_age = round(data["Age"].mean(), 1)

   st.metric(label="📊 Avg Age", value=f"{avg_age} yrs")

# ---------------- □ VISUALIZATION GRID (MIDDLE ROW) -----------------

st.markdown("### 📊 Data Visualizations")

row1_col1, row1_col2 = st.columns(2)

# 📌 Bar Chart - Survival Count
```

```
with row1_col1:
    st.markdown("#### 📌 Survival Count")
    fig, ax = plt.subplots()
    data["Survived"].value_counts().plot(kind="bar", color=["red", "green"], ax=ax)
    plt.xlabel("Survival (0 = No, 1 = Yes)")
    plt.ylabel("Count")
    st.pyplot(fig)
# 📌 Histogram - Age Distribution
with row1_col2:
    st.markdown("#### 📌 Age Distribution of Passengers")
    fig, ax = plt.subplots()
    sns.histplot(data["Age"].dropna(), bins=20, kde=True, ax=ax)
    plt.xlabel("Age")
    plt.ylabel("Frequency")
    st.pyplot(fig)
# ---------- 🔹 INTERACTIVE SCATTER PLOT (FULL WIDTH) ----------
st.markdown("### 🔍 Passenger Insights")
fig = px.scatter(data, x="Age", y="Fare", color="Survived",
                 title="Fare Paid vs Age",
                 labels={"Fare": "Fare Paid", "Age": "Passenger Age"},
                 width=900, height=500)
st.plotly_chart(fig, use_container_width=True)
# ---------- 🔹 FILTERED DATA TABLE (BOTTOM ROW) ----------
st.markdown("### 🎯 Filter & Explore Data")
pclass_filter = st.selectbox("🔍 Select Passenger Class", sorted(data["Pclass"].unique()))
filtered_data = data[data["Pclass"] == pclass_filter]
st.dataframe(filtered_data)
# ---------- 🔹 STREAMLIT DASHBOARD END ----------
```

OUTPUT:-

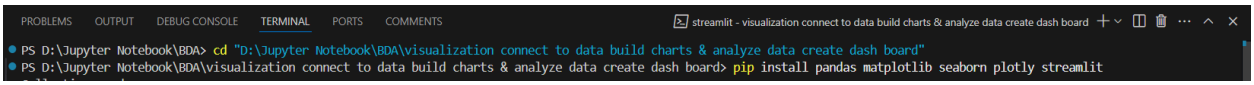

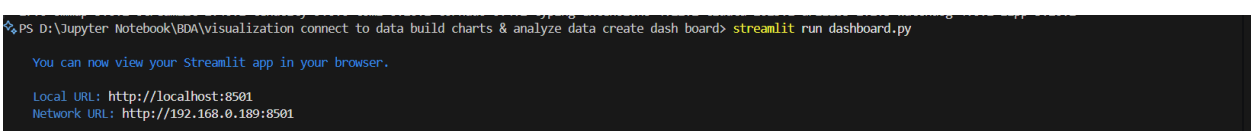

Titanic Passenger Dashboard

Analyze survival rates and passenger demographics

Key Statistics

Total Passengers	Survived	Avg Fare	Avg Age
891	342	$32.2	29.7 yrs

Data Visualizations

Survival Count

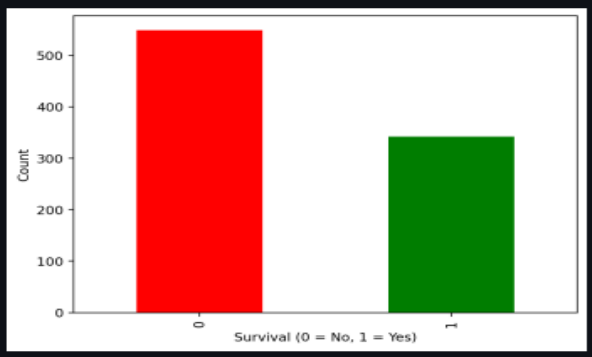

Age Distribution of Passengers

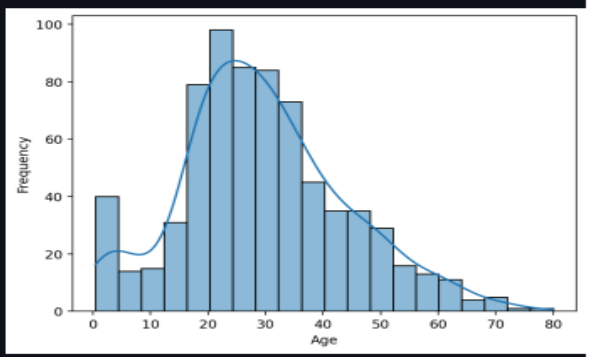

Passenger Insights

Fare Paid vs Age

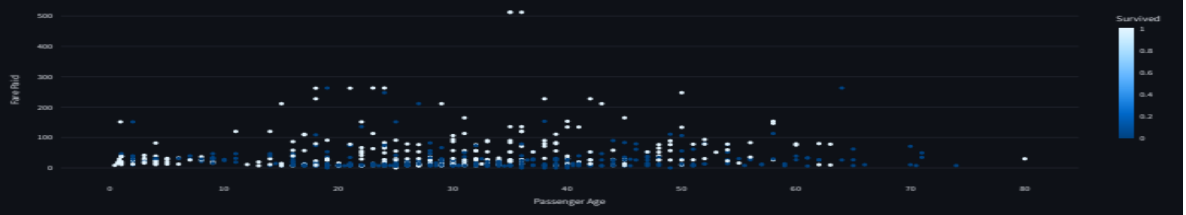

Filter & Explore Data

Select Passenger Class

1

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38	1	0	PC 17599	71.2833	C85	C
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35	1	0	113803	53.1	C123	S
6	7	0	1	McCarthy, Mr. Timothy J	male	54	0	0	17463	51.8625	E46	S
11	12	1	1	Bonnell, Miss. Elizabeth	female	58	0	0	113783	26.55	C103	S
23	24	1	1	Sloper, Mr. William Thompson	male	28	0	0	113788	35.5	A6	S
27	28	0	1	Fortune, Mr. Charles Alexander	male	19	3	2	19950	263	C23 C2	S
30	31	0	1	Uruchurtu, Don. Manuel E	male	40	0	0	PC 17601	27.7208	None	C
31	32	1	1	Spencer, Mrs. William Augustus (Marie Eugenie)	female	None	1	0	PC 17569	146.5208	B78	C
34	35	0	1	Meyer, Mr. Edgar Joseph	male	28	1	0	PC 17604	82.1708	None	C
35	36	0	1	Holverson, Mr. Alexander Oskar	male	42	1	0	113789	52	None	S